

Breast Cancer Prediction with ML

CS-715

Howard Nguyen

22 April, 2023

Contents

1	Methods and Analysis	2
1.1	Data Analysis	2
1.2	Modelling Approach	10
1.2.1	Modelling	10
1.2.2	Model creation	15
1.2.3	Naive Bayes Model	16
1.2.4	Logistic Regression Model	17
1.2.5	Random Forest Model	19
1.2.6	K Nearest Neighbor (KNN) Model	21
1.2.7	Neural Network with PCA Model	23
1.2.8	Neural Network with LDA Model	25
2	Results	27
3	Discussion	30
4	Conclusion	32
5	Appendix - Environment	33

The dataset's features describe characteristics of the cell nuclei on the image. The features information are specified below:

- Attribute Information:
 1. ID number
 2. Diagnosis (M = Malignant, B = Benign)
- Ten features were computed for each cell nucleus:
 1. radius: mean of distances from center to points on the perimeter
 2. texture: standard deviation of grey-scale values
 3. perimeter
 4. area: Number of pixels inside contour + $\frac{1}{2}$ for pixels on perimeter
 5. smoothness: local variation in radius lengths), , t
 6. compactness: $\text{perimeter}^2 / \text{area} - 1.0$; this dimensionless number is at a minimum with a circular disk and increases with the irregularity of the boundary, but this measure also increases for elongated cell nuclei, which is not indicative of malignancy
 7. concavity: severity of concave portions of the contour
 8. concave points: number of concave portions of the contour
 9. symmetry
 10. fractal dimension: "coastline approximation" - 1; a higher value corresponds a less regular contour and thus to a higher probability of malignancy

The mean, standard error and "worst" or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 variables. From this diagnosis, 714 of the cases were classified as benign tumors and 424 were considered malignant tumors. All cancers and some of the benign masses were histologically confirmed

The column 33 is invalid.

```
# View quantity of each value in the diagnosis variable
table(data$diagnosis)
```

```

  B    M
714 424
```

```
# Convert from char to factor for diagnosis variable
data$diagnosis <- as.factor(data$diagnosis)
# the 33 column is invalid
data[,33] <- NULL
```

Chapter 1

Methods and Analysis

1.1 Data Analysis

By observing the dataset, we found that it contains 1138 observations with 32 variables.

```
str(data)
```

```
'data.frame':  1138 obs. of  32 variables:
 $ id                : int  842302 842517 84300903 84348301 84358402 843786 844359 84458202 844981
 $ diagnosis         : Factor w/ 2 levels "B","M": 2 2 2 2 2 2 2 2 2 2 ...
 $ radius_mean       : num  18 20.6 19.7 11.4 20.3 ...
 $ texture_mean      : num  10.4 17.8 21.2 20.4 14.3 ...
 $ perimeter_mean    : num  122.8 132.9 130 77.6 135.1 ...
 $ area_mean         : num  1001 1326 1203 386 1297 ...
 $ smoothness_mean   : num  0.1184 0.0847 0.1096 0.1425 0.1003 ...
 $ compactness_mean  : num  0.2776 0.0786 0.1599 0.2839 0.1328 ...
 $ concavity_mean    : num  0.3001 0.0869 0.1974 0.2414 0.198 ...
 $ concave.points_mean : num  0.1471 0.0702 0.1279 0.1052 0.1043 ...
 $ symmetry_mean     : num  0.242 0.181 0.207 0.26 0.181 ...
 $ fractal_dimension_mean : num  0.0787 0.0567 0.06 0.0974 0.0588 ...
 $ radius_se         : num  1.095 0.543 0.746 0.496 0.757 ...
 $ texture_se        : num  0.905 0.734 0.787 1.156 0.781 ...
 $ perimeter_se      : num  8.59 3.4 4.58 3.44 5.44 ...
 $ area_se           : num  153.4 74.1 94 27.2 94.4 ...
 $ smoothness_se     : num  0.0064 0.00522 0.00615 0.00911 0.01149 ...
 $ compactness_se    : num  0.049 0.0131 0.0401 0.0746 0.0246 ...
 $ concavity_se      : num  0.0537 0.0186 0.0383 0.0566 0.0569 ...
 $ concave.points_se : num  0.0159 0.0134 0.0206 0.0187 0.0188 ...
 $ symmetry_se       : num  0.03 0.0139 0.0225 0.0596 0.0176 ...
 $ fractal_dimension_se : num  0.00619 0.00353 0.00457 0.00921 0.00511 ...
 $ radius_worst      : num  25.4 25 23.6 14.9 22.5 ...
 $ texture_worst     : num  17.3 23.4 25.5 26.5 16.7 ...
 $ perimeter_worst   : num  184.6 158.8 152.5 98.9 152.2 ...
 $ area_worst        : num  2019 1956 1709 568 1575 ...
 $ smoothness_worst  : num  0.162 0.124 0.144 0.21 0.137 ...
 $ compactness_worst : num  0.666 0.187 0.424 0.866 0.205 ...
 $ concavity_worst   : num  0.712 0.242 0.45 0.687 0.4 ...
 $ concave.points_worst : num  0.265 0.186 0.243 0.258 0.163 ...
```

```
$ symmetry_worst      : num  0.46 0.275 0.361 0.664 0.236 ...
$ fractal_dimension_worst: num  0.1189 0.089 0.0876 0.173 0.0768 ...
```

```
head(data)
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
1	842302	M	17.99	10.38	122.80	1001.0
2	842517	M	20.57	17.77	132.90	1326.0
3	84300903	M	19.69	21.25	130.00	1203.0
4	84348301	M	11.42	20.38	77.58	386.1
5	84358402	M	20.29	14.34	135.10	1297.0
6	843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
1	0.11840	0.27760	0.3001	0.14710
2	0.08474	0.07864	0.0869	0.07017
3	0.10960	0.15990	0.1974	0.12790
4	0.14250	0.28390	0.2414	0.10520
5	0.10030	0.13280	0.1980	0.10430
6	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
1	0.2419	0.07871	1.0950	0.9053	8.589
2	0.1812	0.05667	0.5435	0.7339	3.398
3	0.2069	0.05999	0.7456	0.7869	4.585
4	0.2597	0.09744	0.4956	1.1560	3.445
5	0.1809	0.05883	0.7572	0.7813	5.438
6	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
1	153.40	0.006399	0.04904	0.05373	0.01587
2	74.08	0.005225	0.01308	0.01860	0.01340
3	94.03	0.006150	0.04006	0.03832	0.02058
4	27.23	0.009110	0.07458	0.05661	0.01867
5	94.44	0.011490	0.02461	0.05688	0.01885
6	27.19	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	perimeter_worst
1	0.03003	0.006193	25.38	17.33	184.60
2	0.01389	0.003532	24.99	23.41	158.80
3	0.02250	0.004571	23.57	25.53	152.50
4	0.05963	0.009208	14.91	26.50	98.87
5	0.01756	0.005115	22.54	16.67	152.20
6	0.02165	0.005082	15.47	23.75	103.40

	area_worst	smoothness_worst	compactness_worst	concavity_worst
1	2019.0	0.1622	0.6656	0.7119
2	1956.0	0.1238	0.1866	0.2416
3	1709.0	0.1444	0.4245	0.4504
4	567.7	0.2098	0.8663	0.6869
5	1575.0	0.1374	0.2050	0.4000
6	741.6	0.1791	0.5249	0.5355

	concave.points_worst	symmetry_worst	fractal_dimension_worst
1	0.2654	0.4601	0.11890
2	0.1860	0.2750	0.08902
3	0.2430	0.3613	0.08758
4	0.2575	0.6638	0.17300
5	0.1625	0.2364	0.07678
6	0.1741	0.3985	0.12440

```
summary(data)
```

```

      id      diagnosis radius_mean texture_mean
Min.   : 8670      B:714    Min.   : 6.981    Min.   : 9.71
1st Qu.: 869218    M:424    1st Qu.:11.700    1st Qu.:16.17
Median : 906024                Median :13.370    Median :18.84
Mean   : 30371831                Mean   :14.127    Mean   :19.29
3rd Qu.: 8813129                3rd Qu.:15.780    3rd Qu.:21.80
Max.   :911320502                Max.   :28.110    Max.   :39.28

perimeter_mean area_mean smoothness_mean compactness_mean
Min.   : 43.79    Min.   : 143.5    Min.   :0.05263    Min.   :0.01938
1st Qu.: 75.17    1st Qu.: 420.3    1st Qu.:0.08637    1st Qu.:0.06492
Median : 86.24    Median : 551.1    Median :0.09587    Median :0.09263
Mean   : 91.97    Mean   : 654.9    Mean   :0.09636    Mean   :0.10434
3rd Qu.:104.10    3rd Qu.: 782.7    3rd Qu.:0.10530    3rd Qu.:0.13040
Max.   :188.50    Max.   :2501.0    Max.   :0.16340    Max.   :0.34540

concavity_mean concave.points_mean symmetry_mean fractal_dimension_mean
Min.   :0.00000    Min.   :0.00000    Min.   :0.1060    Min.   :0.04996
1st Qu.:0.02956    1st Qu.:0.02031    1st Qu.:0.1619    1st Qu.:0.05770
Median :0.06154    Median :0.03350    Median :0.1792    Median :0.06154
Mean   :0.08880    Mean   :0.04892    Mean   :0.1812    Mean   :0.06280
3rd Qu.:0.13070    3rd Qu.:0.07400    3rd Qu.:0.1957    3rd Qu.:0.06612
Max.   :0.42680    Max.   :0.20120    Max.   :0.3040    Max.   :0.09744

radius_se texture_se perimeter_se area_se
Min.   :0.1115    Min.   :0.3602    Min.   : 0.757    Min.   : 6.802
1st Qu.:0.2324    1st Qu.:0.8339    1st Qu.: 1.606    1st Qu.:17.850
Median :0.3242    Median :1.1080    Median : 2.287    Median :24.530
Mean   :0.4052    Mean   :1.2169    Mean   : 2.866    Mean   :40.337
3rd Qu.:0.4789    3rd Qu.:1.4740    3rd Qu.: 3.357    3rd Qu.:45.190
Max.   :2.8730    Max.   :4.8850    Max.   :21.980    Max.   :542.200

smoothness_se compactness_se concavity_se concave.points_se
Min.   :0.001713    Min.   :0.002252    Min.   :0.00000    Min.   :0.000000
1st Qu.:0.005169    1st Qu.:0.013080    1st Qu.:0.01509    1st Qu.:0.007638
Median :0.006380    Median :0.020450    Median :0.02589    Median :0.010930
Mean   :0.007041    Mean   :0.025478    Mean   :0.03189    Mean   :0.011796
3rd Qu.:0.008146    3rd Qu.:0.032450    3rd Qu.:0.04205    3rd Qu.:0.014710
Max.   :0.031130    Max.   :0.135400    Max.   :0.39600    Max.   :0.052790

symmetry_se fractal_dimension_se radius_worst texture_worst
Min.   :0.007882    Min.   :0.0008948    Min.   : 7.93    Min.   :12.02
1st Qu.:0.015160    1st Qu.:0.0022480    1st Qu.:13.01    1st Qu.:21.08
Median :0.018730    Median :0.0031870    Median :14.97    Median :25.41
Mean   :0.020542    Mean   :0.0037949    Mean   :16.27    Mean   :25.68
3rd Qu.:0.023480    3rd Qu.:0.0045580    3rd Qu.:18.79    3rd Qu.:29.72
Max.   :0.078950    Max.   :0.0298400    Max.   :36.04    Max.   :49.54

perimeter_worst area_worst smoothness_worst compactness_worst
Min.   : 50.41    Min.   : 185.2    Min.   :0.07117    Min.   :0.02729
1st Qu.: 84.11    1st Qu.: 515.3    1st Qu.:0.11660    1st Qu.:0.14720
Median : 97.66    Median : 686.5    Median :0.13130    Median :0.21190
Mean   :107.26    Mean   : 880.6    Mean   :0.13237    Mean   :0.25427
3rd Qu.:125.40    3rd Qu.:1084.0    3rd Qu.:0.14600    3rd Qu.:0.33910
Max.   :251.20    Max.   :4254.0    Max.   :0.22260    Max.   :1.05800

concavity_worst concave.points_worst symmetry_worst fractal_dimension_worst
Min.   :0.0000    Min.   :0.00000    Min.   :0.1565    Min.   :0.05504

```

1st Qu.:0.1145	1st Qu.:0.06493	1st Qu.:0.2504	1st Qu.:0.07146
Median :0.2267	Median :0.09993	Median :0.2822	Median :0.08004
Mean :0.2722	Mean :0.11461	Mean :0.2901	Mean :0.08395
3rd Qu.:0.3829	3rd Qu.:0.16140	3rd Qu.:0.3179	3rd Qu.:0.09208
Max. :1.2520	Max. :0.29100	Max. :0.6638	Max. :0.20750

We have to check if the dataset has any missing value:

```
$id
[1] 0
```

```
$diagnosis
[1] 0
```

```
$radius_mean
[1] 0
```

```
$texture_mean
[1] 0
```

```
$perimeter_mean
[1] 0
```

```
$area_mean
[1] 0
```

```
$smoothness_mean
[1] 0
```

```
$compactness_mean
[1] 0
```

```
$concavity_mean
[1] 0
```

```
$concave.points_mean
[1] 0
```

```
$symmetry_mean
[1] 0
```

```
$fractal_dimension_mean
[1] 0
```

```
$radius_se
[1] 0
```

```
$texture_se
[1] 0
```

```
$perimeter_se
[1] 0
```

```
$area_se
```

```
[1] 0

$smoothness_se
[1] 0

$compactness_se
[1] 0

$concavity_se
[1] 0

$concave.points_se
[1] 0

$symmetry_se
[1] 0

$fractal_dimension_se
[1] 0

$radius_worst
[1] 0

$texture_worst
[1] 0

$perimeter_worst
[1] 0

$area_worst
[1] 0

$smoothness_worst
[1] 0

$compactness_worst
[1] 0

$concavity_worst
[1] 0

$concave.points_worst
[1] 0

$symmetry_worst
[1] 0

$fractal_dimension_worst
[1] 0
```


It results that there aren't NA values. By analyzing the the dataset we discover that it is a bit unbalanced in its proportions:


```
prop.table(table(data$diagnosis))
```

```
      B      M
0.6274165 0.3725835
```

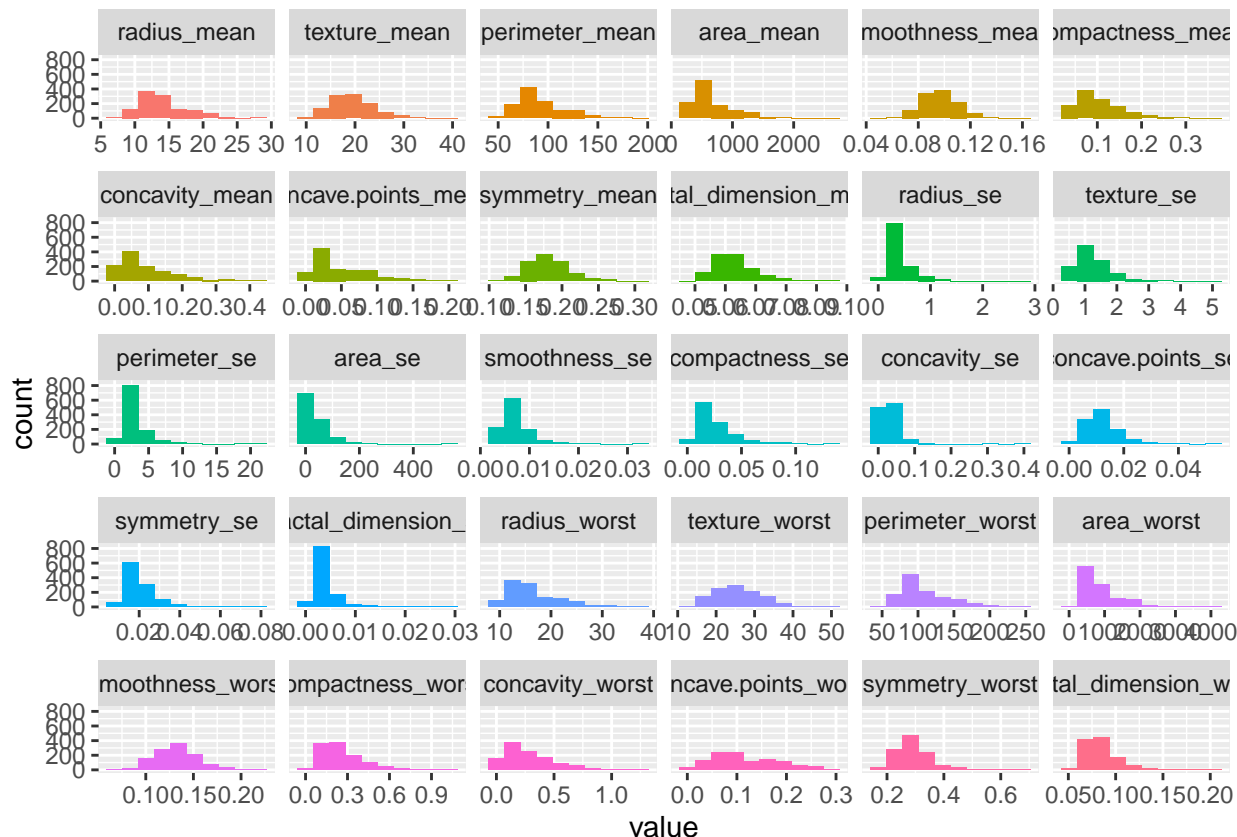
Also the plot of proportions confirms that the target variable is slightly unbalanced.

```
options(repr.plot.width=4, repr.plot.height=4)
ggplot(data, aes(x=diagnosis)) +
  geom_bar(fill="sky blue", alpha=0.5) +
  theme_bw() +
  labs(title="Distribution of Diagnosis")
```




The most variables of the dataset are normally distributed as show with the below plot:

```
plot_num(data %>% select(-id), bins=10)
```



Now we have to check if there is any correlation between variables as machine learning algorithms assume that the predictor variables are independent from each others.

```
correlationMatrix <- cor(data[,3:ncol(data)])
corrplot(correlationMatrix,
  order = "hclust", # order for labels, can be "original"
  type = "upper",   # matrix: full, upper, lower
  diag = F,        # remove diagonal
  tl.cex = 0.6,    # font size
  tl.srt = 75,     # label angle
  tl.col = "black",
  addrect = 8)
```



As shown by this plot, many variables are highly correlated with each others. Many methods perform better if highly correlated attributes are removed. The Caret R package provides the `findCorrelation` which will analyze a correlation matrix of your data's attributes report on attributes that can be removed. Because of many correlations may cause some of machine learning models fail.

```
# find attributes that are highly corrected (ideally >0.90)
highlyCorrelated <- findCorrelation(correlationMatrix, cutoff=0.9)
# print indexes of highly correlated attributes
print(highlyCorrelated)
```

```
[1] 7 8 23 21 3 24 1 13 14 2
```

Selecting the right features in our data can mean the difference between mediocre performance with long training times and great performance with short training times.

```
# Remove correlated variables
data2 <- data %>%select(-highlyCorrelated)
# number of columns after removing correlated variables
ncol(data2)
```

```
[1] 22
```

The new dataset has been removed 10 variables, it's 22 variables now.


1.2 Modelling Approach

1.2.1 Modelling

Principal Component Analysis (PCA).

To avoid redundancy and relevancy, we used the function ‘prcomp’ to calculate the Principal Component Analysis (PCA) and select the rights components to avoid correlated variables that can be detrimental to our clustering analysis. One of the common problems in analysis of complex data comes from a large number of variables, which requires a large amount of memory and computation power. This is where PCA comes in. It is a technique to reduce the dimension of the feature space by feature extraction. The main idea of PCA is to reduce the dimension of a data set consisting of many variables correlated with each other, either heavily or lightly, while retaining the variation present in the dataset, up to the maximum extent. The same is done by transforming the variables to a new set of variables, which are known as the principal components (or simply, the PCs) and are orthogonal, ordered such that the retention of variation present in the original variables decrease as we move down in the order.

```
pca_res_data <- prcomp(data[,3:ncol(data)], center = TRUE, scale = TRUE)
plot(pca_res_data, type="l")
```



```
summary(pca_res_data)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172

Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

As we can observe from the above table, the two first components explains the 0.6324 of the variance. We need 10 principal components to explain more than 0.95 of the variance and 17 to explain more than 0.99.

```
pca_res_data2 <- prcomp(data2[,3:ncol(data2)], center = TRUE, scale = TRUE)
plot(pca_res_data2, type="l")
```




```
summary(pca_res_data2)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.0980	2.0196	1.4663	1.0845	0.91561	0.77019	0.57227
Proportion of Variance	0.4799	0.2039	0.1075	0.0588	0.04192	0.02966	0.01637
Cumulative Proportion	0.4799	0.6838	0.7913	0.8501	0.89205	0.92171	0.93808
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.53641	0.50898	0.45726	0.36641	0.31778	0.28802	0.21369
Proportion of Variance	0.01439	0.01295	0.01045	0.00671	0.00505	0.00415	0.00228
Cumulative Proportion	0.95247	0.96542	0.97588	0.98259	0.98764	0.99179	0.99407
	PC15	PC16	PC17	PC18	PC19	PC20	
Standard deviation	0.1846	0.15579	0.15393	0.14782	0.09636	0.07375	
Proportion of Variance	0.0017	0.00121	0.00118	0.00109	0.00046	0.00027	
Cumulative Proportion	0.9958	0.99699	0.99817	0.99926	0.99973	1.00000	


The above table shows that 95% of the variance is explained with 8 PC's in the transformed dataset data2.

```
pca_df <- as.data.frame(pca_res_data2$x)
ggplot(pca_df, aes(x=PC1, y=PC2, col=data$diagnosis)) + geom_point(alpha=0.5)
```



The data of the first 2 components can be easily separated into two classes. This is caused by the fact that the variance explained by these components is not large. The data can be easily separated.

```
g_pc1 <- ggplot(pca_df, aes(x=PC1, fill=data$diagnosis)) + geom_density(alpha=0.25)
g_pc2 <- ggplot(pca_df, aes(x=PC2, fill=data$diagnosis)) + geom_density(alpha=0.25)
grid.arrange(g_pc1, g_pc2, ncol=2)
```



Linear Discriminant Analysis (LDA)

Another approach is to use the Linear Discriminant Analysis (LDA) instead of PCA. LDA takes in consideration the different classes and could get better results. The particularity of LDA is that it models the distribution of predictors separately in each of the response classes, and then it uses Bayes' theorem to estimate the probability. It is important to know that LDA assumes a normal distribution for each class, a class-specific mean, and a common variance.

```
lda_res_data <- MASS::lda(diagnosis~., data = data, center = TRUE, scale = TRUE)
lda_res_data
```

Call:

```
lda(diagnosis ~ ., data = data, center = TRUE, scale = TRUE)
```

Prior probabilities of groups:

	B	M
	0.6274165	0.3725835

Group means:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
B	26543825	12.14652	17.91476	78.07541	462.7902	0.09247765
M	36818050	17.46283	21.60491	115.36538	978.3764	0.10289849

	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
B	0.08008462	0.04605762	0.02571741	0.174186	
M	0.14518778	0.16077472	0.08799000	0.192909	
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
B	0.06286739	0.2840824	1.220380	2.000321	21.13515
M	0.06268009	0.6090825	1.210915	4.323929	72.67241
	smoothness_se	compactness_se	concavity_se	concave.points_se	symmetry_se
B	0.007195902	0.02143825	0.02599674	0.009857653	0.02058381
M	0.006780094	0.03228117	0.04182401	0.015060472	0.02047240
	fractal_dimension_se	radius_worst	texture_worst	perimeter_worst	area_worst
B	0.003636051	13.37980	23.51507	87.00594	558.8994
M	0.004062406	21.13481	29.31821	141.37033	1422.2863
	smoothness_worst	compactness_worst	concavity_worst	concave.points_worst	
B	0.1249595	0.1826725	0.1662377	0.07444434	
M	0.1448452	0.3748241	0.4506056	0.18223731	
	symmetry_worst	fractal_dimension_worst			
B	0.2702459	0.07944207			
M	0.3234679	0.09152995			


Coefficients of linear discriminants:

	LD1
id	-2.514331e-10
radius_mean	-1.081828e+00
texture_mean	2.340469e-02
perimeter_mean	1.173741e-01
area_mean	1.597096e-03
smoothness_mean	5.256204e-01
compactness_mean	-2.096043e+01
concavity_mean	6.962054e+00
concave.points_mean	1.048491e+01
symmetry_mean	4.943251e-01
fractal_dimension_mean	-5.942897e-02
radius_se	2.103355e+00
texture_se	-3.983377e-02
perimeter_se	-1.122803e-01
area_se	-4.087103e-03
smoothness_se	7.994704e+01
compactness_se	1.388249e-01
concavity_se	-1.769819e+01
concave.points_se	5.355236e+01
symmetry_se	8.150789e+00
fractal_dimension_se	-3.434380e+01
radius_worst	9.685737e-01
texture_worst	3.543712e-02
perimeter_worst	-1.205569e-02
area_worst	-5.016545e-03
smoothness_worst	2.614560e+00
compactness_worst	3.640098e-01
concavity_worst	1.882357e+00
concave.points_worst	2.220144e+00
symmetry_worst	2.785556e+00
fractal_dimension_worst	2.119696e+01


```
#Data frame of the LDA for visualization purposes
```

```
lda_df_predict <-  
  predict(lda_res_data, data)$x %>%  
  as.data.frame() %>%  
  cbind(diagnosis=data$diagnosis)
```

```
ggplot(lda_df_predict, aes(x=LD1, fill=diagnosis)) +  
  geom_density(alpha=0.5)
```



1.2.2 Model creation

We are going to get a training and a testing set to use when building some models. We split the modified dataset into Train (80%) and Test (20%), in order to predict whether it is a cancer cell is Benign or Malignant, by building machine learning classification models.

```
set.seed(1815)  
data3 <- cbind (diagnosis=data$diagnosis, data2)  
data_sampling_index <- createDataPartition(data$diagnosis, times=1, p=0.8, list = FALSE)  
train_data <- data3[data_sampling_index, ]  
test_data <- data3[-data_sampling_index, ]
```

```
fitControl <- trainControl(method="cv",  
  #Control the computational nuances of the train function)
```

```

number = 15,
#Either the number of folds or number of resampling iterations
classProbs = TRUE,
summaryFunction = twoClassSummary)

```

1.2.3 Naive Bayes Model

The Naive Bayesian classifier is based on Bayes' theorem with the independence assumptions between predictors. A Naive Bayesian model is easy to build, with no complicated iterative parameter estimation which makes it particularly useful for very large datasets. Bayes theorem provides a way of calculating the posterior probability, $P(c|x)$, from $P(c)$, $P(x)$, and $P(x|c)$. Naive Bayes classifier assume that the effect of the value of a predictor (x) on a given class (c) is independent of the values of other predictors. This assumption is called class conditional independence.

```

model_naiveb <- train(diagnosis~.,
                      train_data,
                      method="nb",
                      metric="ROC",
                      preProcess=c('center', 'scale'), #in order to normalize the data
                      trace=FALSE,
                      trControl=fitControl)

prediction_naiveb <- predict(model_naiveb, test_data)
confusionmatrix_naiveb <- confusionMatrix(prediction_naiveb, test_data$diagnosis,
                                           positive = "M")

confusionmatrix_naiveb

```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	131	8
M	11	76

```

Accuracy : 0.9159
95% CI : (0.8718, 0.9486)
No Information Rate : 0.6283
P-Value [Acc > NIR] : <2e-16

```

```
Kappa : 0.8213
```

```
McNemar's Test P-Value : 0.6464
```

```

Sensitivity : 0.9048
Specificity : 0.9225
Pos Pred Value : 0.8736
Neg Pred Value : 0.9424
Prevalence : 0.3717
Detection Rate : 0.3363
Detection Prevalence : 0.3850
Balanced Accuracy : 0.9136

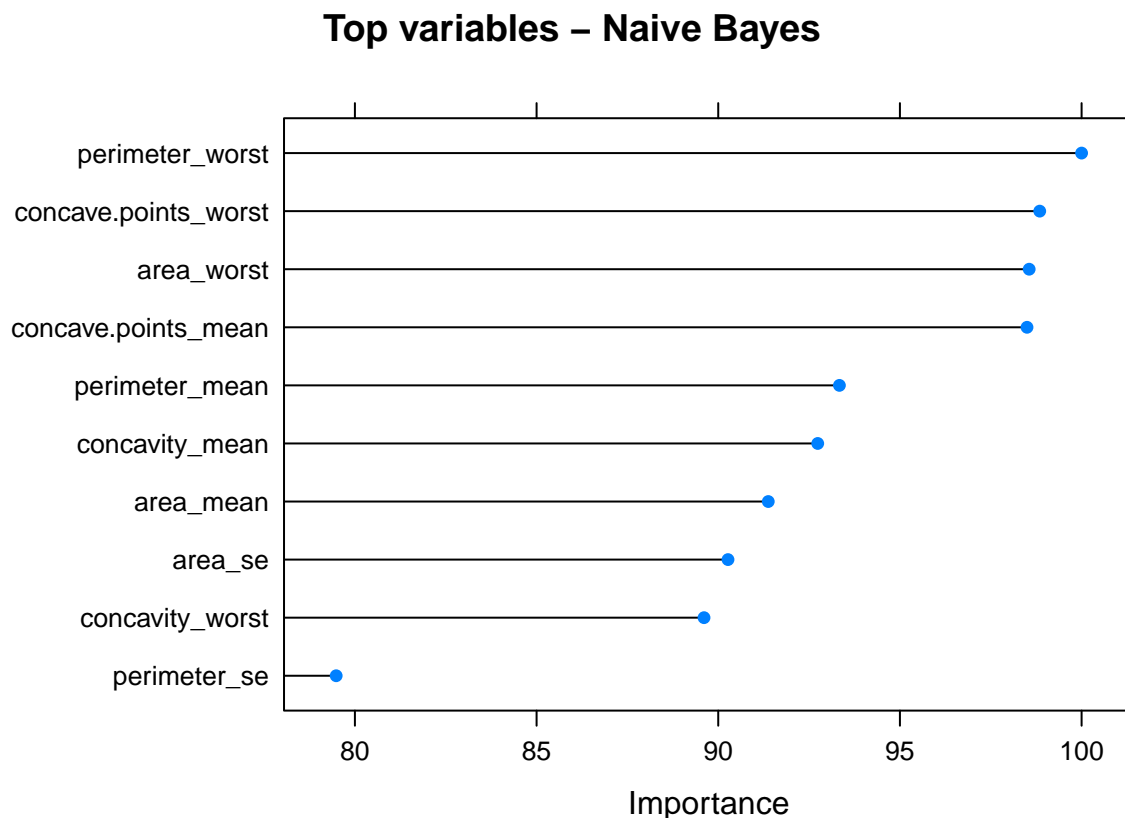
```

'Positive' Class : M

We can note the accuracy with such model. We will later describe better these metrics, where: Sensitivity (recall) represent the true positive rate: the proportions of actual positives correctly identified. Specificity is the true negative rate: the proportion of actual negatives correctly identified. Accuracy is the general score of the classifier model performance as it is the ratio of how many samples are correctly classified to all samples. F1 score: the harmonic mean of precision and sensitivity. Accuracy and F1 score would be used to compare the result with the benchmark model. Precision: the number of correct positive results divided by the number of all positive results returned by the classifier.

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_naiveb), top=10, main="Top variables - Naive Bayes")
```



1.2.4 Logistic Regression Model

Logistic Regression is widely used for binary classification like (0,1). The binary logistic model is used to estimate the probability of a binary response based on one or more predictor (or independent) variables (features).

```
model_logreg<- train(diagnosis ~., data = train_data, method = "glm",
                      metric = "ROC",
```

```

preProcess = c("scale", "center"), # in order to normalize the data
trControl= fitControl)
prediction_logreg<- predict(model_logreg, test_data)

# Check results
confusionmatrix_logreg <- confusionMatrix(prediction_logreg, test_data$diagnosis,
                                           positive = "M")
confusionmatrix_logreg

```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	138	4
M	4	80

Accuracy : 0.9646
 95% CI : (0.9314, 0.9846)
 No Information Rate : 0.6283
 P-Value [Acc > NIR] : <2e-16

 Kappa : 0.9242


 McNemar's Test P-Value : 1

 Sensitivity : 0.9524
 Specificity : 0.9718
 Pos Pred Value : 0.9524
 Neg Pred Value : 0.9718
 Prevalence : 0.3717
 Detection Rate : 0.3540
 Detection Prevalence : 0.3717
 Balanced Accuracy : 0.9621

 'Positive' Class : M

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_logreg), top=10, main="Top variables - Log Regr")
```



1.2.5 Random Forest Model

Random forests are a very popular machine learning approach that addresses the shortcomings of decision trees using a clever idea. The goal is to improve prediction performance and reduce instability by averaging multiple decision trees (a forest of trees constructed with randomness). Random forest is another ensemble method based on decision trees. It splits data into sub-samples, trains decision tree classifiers on each sub-sample and averages prediction of each classifier. Splitting dataset causes higher bias but it is compensated by large decrease in variance. Random Forest is a supervised learning algorithm and it is flexible, easy to use machine learning algorithm that produces, even without hyper-parameter tuning, a great result most of the time. It is also one of the most used algorithms, because of its simplicity and the fact that it can be used for both classification and regression tasks. Random forest builds multiple decision trees and merges them together to get a more accurate and stable prediction.

```
model_randomforest <- train(diagnosis~.,
                             train_data,
                             method="rf", # also recommended ranger, because it is
                                           # a lot faster than original RandomForest (rf)
                             metric="ROC",
                             #tuneLength=10,
                             #tuneGrid = expand.grid(mtry = c(2, 3, 6)),
                             preProcess = c('center', 'scale'),
                             trControl=fitControl)

prediction_randomforest <- predict(model_randomforest, test_data)
```

```
#Check results
confusionmatrix_randomforest <- confusionMatrix(prediction_randomforest,
                                                    test_data$diagnosis, positive = "M")
confusionmatrix_randomforest
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	136	0
M	6	84

Accuracy : 0.9735
 95% CI : (0.9431, 0.9902)
 No Information Rate : 0.6283
 P-Value [Acc > NIR] : < 2e-16

Kappa : 0.944


Mcnemar's Test P-Value : 0.04123

Sensitivity : 1.0000
 Specificity : 0.9577
 Pos Pred Value : 0.9333
 Neg Pred Value : 1.0000
 Prevalence : 0.3717
 Detection Rate : 0.3717
 Detection Prevalence : 0.3982
 Balanced Accuracy : 0.9789

'Positive' Class : M

```
plot(varImp(model_randomforest), top=10, main="Top variables - Random Forest")
```

Top variables – Random Forest



1.2.6 K Nearest Neighbor (KNN) Model

KNN (K-Nearest Neighbors) is one of many (supervised learning) algorithms used in data mining and machine learning, it's a classifier algorithm where the learning is based "how similar" is a data from other. K nearest neighbors is a simple algorithm that stores all available cases and classifies new cases based on a similarity measure (e.g., distance functions).

```
model_knn <- train(diagnosis~.,
  train_data,
  method="knn",
  metric="ROC",
  preProcess = c('center', 'scale'),
  tuneLength=10, #The tune Length parameter tells the algorithm
                 # to try different default values for the main
                 # parameter
                 # In this case we used 10 default values
  trControl=fitControl)

prediction_knn <- predict(model_knn, test_data)
confusionmatrix_knn <- confusionMatrix(prediction_knn, test_data$diagnosis, positive = "M")
confusionmatrix_knn
```

Confusion Matrix and Statistics

```

      Reference
Prediction  B   M
      B 139   6
      M   3  78

      Accuracy : 0.9602
      95% CI : (0.9258, 0.9816)
      No Information Rate : 0.6283
      P-Value [Acc > NIR] : <2e-16

      Kappa : 0.9141

      McNemar's Test P-Value : 0.505


      Sensitivity : 0.9286
      Specificity : 0.9789
      Pos Pred Value : 0.9630
      Neg Pred Value : 0.9586
      Prevalence : 0.3717
      Detection Rate : 0.3451
      Detection Prevalence : 0.3584
      Balanced Accuracy : 0.9537

      'Positive' Class : M

```

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_knn), top=10, main="Top variables - KNN")
```

1.2.7 Neural Network with PCA Model

Artificial Neural Networks (NN) are a types of mathematical algorithms originating in the simulation of networks of biological neurons. An artificial Neural Network consists of nodes (called neurons) and edges (called synapses). Input data is transmitted through the weighted synapses to the neurons where calculations are processed and then either sent to further neurons or represent the output.

Neural Networks take in the weights of connections between neurons . The weights are balanced, learning data point in the wake of learning data point. When all weights are trained, the neural network can be utilized to predict the class or a quantity, if there should arise an occurrence of regression of a new input data point. With Neural networks, extremely complex models can be trained and they can be utilized as a kind of black box, without playing out an unpredictable complex feature engineering before training the model. Joined with the “deep approach” even more unpredictable models can be picked up to realize new possibilities.

```
model_nnet_pca <- train(diagnosis~.,
                        train_data,
                        method="nnet",
                        metric="ROC",
                        preProcess=c('center', 'scale', 'pca'),
                        tuneLength=10,
                        trace=FALSE,
                        trControl=fitControl)

prediction_nnet_pca <- predict(model_nnet_pca, test_data)
```

```
confusionmatrix_nnet_pca <- confusionMatrix(prediction_nnet_pca,
                                             test_data$diagnosis, positive = "M")
confusionmatrix_nnet_pca
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	142	2
M	0	82

Accuracy : 0.9912
 95% CI : (0.9684, 0.9989)
 No Information Rate : 0.6283
 P-Value [Acc > NIR] : <2e-16

 Kappa : 0.981

 McNemar's Test P-Value : 0.4795


 Sensitivity : 0.9762
 Specificity : 1.0000
 Pos Pred Value : 1.0000
 Neg Pred Value : 0.9861
 Prevalence : 0.3717
 Detection Rate : 0.3628
 Detection Prevalence : 0.3628
 Balanced Accuracy : 0.9881

 'Positive' Class : M

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_nnet_pca), top=8, main="Top variables - NNET PCA")
```

Top variables – NNET PCA



1.2.8 Neural Network with LDA Model

We are going to create a training and test set of LDA data created in previous chapters:

```
train_data_lda <- lda_df_predict[data_sampling_index, ]
test_data_lda <- lda_df_predict[-data_sampling_index, ]

model_nnet_lda <- train(diagnosis~.,
                        train_data_lda,
                        method="nnet",
                        metric="ROC",
                        preProcess=c('center', 'scale'),
                        tuneLength=10,
                        trace=FALSE,
                        trControl=fitControl)

prediction_nnet_lda <- predict(model_nnet_lda, test_data_lda)
confusionmatrix_nnet_lda <- confusionMatrix(prediction_nnet_lda,
                                             test_data_lda$diagnosis, positive = "M")

confusionmatrix_nnet_lda
```

Confusion Matrix and Statistics

Reference

Prediction	B	M
B	139	6
M	3	78

Accuracy : 0.9602

95% CI : (0.9258, 0.9816)

No Information Rate : 0.6283

P-Value [Acc > NIR] : <2e-16

Kappa : 0.9141

Mcnemar's Test P-Value : 0.505

Sensitivity : 0.9286

Specificity : 0.9789

Pos Pred Value : 0.9630

Neg Pred Value : 0.9586

Prevalence : 0.3717

Detection Rate : 0.3451

Detection Prevalence : 0.3584

Balanced Accuracy : 0.9537

'Positive' Class : M

Chapter 2

Results

We can now compare and evaluate the results obtained with the above calculations.

```
models_list <- list(Naive_Bayes=model_naiveb,  
                    Logistic_regr=model_logreg,  
                    Random_Forest=model_randomforest,  
                    KNN=model_knn,  
                    Neural_PCA=model_nnet_pca,  
                    Neural_LDA=model_nnet_lda)  
models_results <- resamples(models_list)  
  
summary(models_results)
```

Call:

```
summary.resamples(object = models_results)
```

Models: Naive_Bayes, Logistic_regr, Random_Forest, KNN, Neural_PCA, Neural_LDA
Number of resamples: 15

ROC

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.9712919	0.9823955	0.9931350	0.9893357	0.9959761	1	0
Logistic_regr	0.9533493	0.9965675	0.9988558	0.9948884	1.0000000	1	0
Random_Forest	0.9838350	1.0000000	1.0000000	0.9986138	1.0000000	1	0
KNN	0.9645309	0.9949163	0.9971396	0.9949567	1.0000000	1	0
Neural_PCA	0.9942792	1.0000000	1.0000000	0.9996186	1.0000000	1	0
Neural_LDA	0.9904306	0.9941492	0.9988558	0.9970182	1.0000000	1	0

Sens


	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.8947368	0.9342105	0.9736842	0.9544309	0.9740216	1	0
Logistic_regr	0.9473684	0.9736842	1.0000000	0.9877643	1.0000000	1	0
Random_Forest	0.9736842	0.9868421	1.0000000	0.9929825	1.0000000	1	0
KNN	0.9736842	1.0000000	1.0000000	0.9947818	1.0000000	1	0
Neural_PCA	0.9736842	1.0000000	1.0000000	0.9982456	1.0000000	1	0
Neural_LDA	0.9473684	0.9736842	1.0000000	0.9877193	1.0000000	1	0

Spec

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.7272727	0.8636364	0.9130435	0.9022398	0.9565217	1	0
Logistic_regr	0.9090909	0.9555336	1.0000000	0.9703557	1.0000000	1	0
Random_Forest	0.9090909	0.9565217	1.0000000	0.9822134	1.0000000	1	0
KNN	0.8260870	0.9130435	0.9545455	0.9413702	0.9565217	1	0
Neural_PCA	0.9565217	1.0000000	1.0000000	0.9913043	1.0000000	1	0
Neural_LDA	0.9090909	0.9555336	0.9565217	0.9642951	1.0000000	1	0

As we can observe from the following plot, four models, Naive_Bayes, KNN, Logistic_regr, and Neural_LDA have great variability, depending of the processed sample :

```
bwplot(models_results, metric="ROC")
```



The Neural Network LDA model achieve a great auc (Area Under the ROC Curve) with some variability. The ROC (Receiver Operating characteristic Curve is a graph showing the performance of a classification model at all classification thresholds) metric measure the auc of the roc curve of each model. This metric is independent of any threshold. Let's remember how these models result with the testing dataset. Prediction classes are obtained by default with a threshold of 0.5 which could not be the best with an unbalanced dataset like this.

```
confusionmatrix_list <- list(
  Naive_Bayes=confusionmatrix_naiveb,
  Logistic_regr=confusionmatrix_logreg,
  Random_Forest=confusionmatrix_randomforest,
  KNN=confusionmatrix_knn,
```

```

Neural_PCA=confusionmatrix_nnet_pca,
Neural_LDA=confusionmatrix_nnet_lda)
confusionmatrix_list_results <- sapply(confusionmatrix_list, function(x) x$byClass)
confusionmatrix_list_results %>% knitr::kable()

```

	Naive_Bayes	Logistic_regr	Random_Forest	KNN	Neural_PCA	Neural_LDA
Sensitivity	0.9047619	0.9523810	1.0000000	0.9285714	0.9761905	0.9285714
Specificity	0.9225352	0.9718310	0.9577465	0.9788732	1.0000000	0.9788732
Pos Pred Value	0.8735632	0.9523810	0.9333333	0.9629630	1.0000000	0.9629630
Neg Pred Value	0.9424460	0.9718310	1.0000000	0.9586207	0.9861111	0.9586207
Precision	0.8735632	0.9523810	0.9333333	0.9629630	1.0000000	0.9629630
Recall	0.9047619	0.9523810	1.0000000	0.9285714	0.9761905	0.9285714
F1	0.8888889	0.9523810	0.9655172	0.9454545	0.9879518	0.9454545
Prevalence	0.3716814	0.3716814	0.3716814	0.3716814	0.3716814	0.3716814
Detection Rate	0.3362832	0.3539823	0.3716814	0.3451327	0.3628319	0.3451327
Detection	0.3849558	0.3716814	0.3982301	0.3584071	0.3628319	0.3584071
Prevalence						
Balanced Accuracy	0.9136486	0.9621060	0.9788732	0.9537223	0.9880952	0.9537223

Chapter 3

Discussion

We will now describe the metrics that we will compare in this section.

Accuracy is our starting point. It is the number of correct predictions made divided by the total number of predictions made, multiplied by 100 to turn it into a percentage.

Precision is the number of True Positives divided by the number of True Positives and False Positives. Put another way, it is the number of positive predictions divided by the total number of positive class values predicted. It is also called the Positive Predictive Value (PPV). A low precision can also indicate a large number of False Positives.

Recall (Sensitivity) is the number of True Positives divided by the number of True Positives and the number of False Negatives. Put another way it is the number of positive predictions divided by the number of positive class values in the test data. It is also called Sensitivity or the True Positive Rate. Recall can be thought of as a measure of a classifiers completeness. A low recall indicates many False Negatives.

The F1 Score is the $2 \times ((\text{precision} \times \text{recall}) / (\text{precision} + \text{recall}))$. It is also called the F Score or the F Measure. Put another way, the F1 score conveys the balance between the precision and the recall.

The best results for sensitivity (detection of breast cancer malign cases) is Neural Network with PCA model which also has a great F1 score.

```
confusionmatrix_results_max <- apply(confusionmatrix_list_results, 1, which.is.max)

output_report <-
  data.frame(metric=names(confusionmatrix_results_max),
             best_model=colnames(confusionmatrix_list_results)[confusionmatrix_results_max],
             value=mapply(function(x,y) {confusionmatrix_list_results[x,y]},
                           names(confusionmatrix_results_max),
                           confusionmatrix_results_max))
rownames(output_report) <- NULL
output_report
```

	metric	best_model	value
1	Sensitivity	Random_Forest	1.0000000
2	Specificity	Neural_PCA	1.0000000
3	Pos Pred Value	Neural_PCA	1.0000000
4	Neg Pred Value	Random_Forest	1.0000000
5	Precision	Neural_PCA	1.0000000
6	Recall	Random_Forest	1.0000000
7	F1	Neural_PCA	0.9879518

8	Prevalence	Neural_LDA	0.3716814
9	Detection Rate	Random_Forest	0.3716814
10	Detection Prevalence	Random_Forest	0.3982301
11	Balanced Accuracy	Neural_PCA	0.9880952

Chapter 4

Conclusion

This research treats the Breast Cancer diagnosis problem as a pattern classification problem. In this report we investigated several machine learning model and we selected the optimal model by selecting a high accuracy level combined with a low rate of false-negatives (the means that the metric is high sensitivity).

The Neural Network with PCA model had the optimal results for F1 (0.9879518), Sensitivity (0.9761905) and Balanced Accuracy (0.9880952)

Chapter 5

Appendix - Environment

```
print("Operating System:")
```

```
[1] "Operating System:"
```

```
version
```

```
platform      -  
arch          aarch64-apple-darwin20  
arch          aarch64  
os            darwin20  
system        aarch64, darwin20  
status  
major         4  
minor         2.2  
year          2022  
month         10  
day           31  
svn rev       83211  
language      R  
version.string R version 4.2.2 (2022-10-31)  
nickname      Innocent and Trusting
```